

# SORTS Tech Report

Sam Wintermute, Joseph Xu, James Irizarry

January 16, 2007

## 1 Overview

The goal of this project is to interface Soar to a real-time strategy (RTS) game. RTS games, such as StarCraft, WarCraft, and Command and Conquer, are multi-player strategy games where a player needs to handle many different tasks- planning base layouts, managing economies, organizing attacks, responding to enemy attacks, and even diplomacy. This is all occurring in real time (as the name implies), presenting an extremely rich, challenging environment for an AI system. Making this environment accessible to Soar provides many opportunities for both utilizing existing capabilities and the development of new capabilities.

### 1.1 ORTS Overview

The Open Real Time Strategy software is a highly configurable game engine used to play real time strategy (RTS) games [2]. The main purpose for ORTS is to serve as an open source, open interface RTS game engine for RTS AI tournaments. ORTS is undergoing active development as of July 2006 at the University of Alberta under the direction of Michael Buro.

There are several reasons why ORTS is especially suitable for use in AI tournaments. It has a (relatively) straightforward C++ API, making interfacing with your favorite AI system easy. All the specific game mechanics, ranging from types of units, actions, and physics, are specified via C++ style scripts called blueprints. This means that ORTS can be easily configured to simulate a wide range of environments, from arbitrarily simple ones like Wumpus World to complex ones like Starcraft. Finally, ORTS has a client/server architecture in which the server maintains the state of the world and only report to the clients information they are supposed to have for a fair game. This is in contrast to most commercial RTS games, in which each client maintains the entire world state and prevents the player from accessing forbidden information such as other players' locations only by hiding them from the GUI. The result is that ORTS is impervious to "memory hack" cheats that are widespread in commercial RTS games. This feature is particularly important if tournaments are to be run across the Internet.

### 1.2 SORTS Overview

SORTS is our middleware system that allows Soar to act as a client to the ORTS game server, so that ORTS game playing agents can be written in and executed on Soar. SORTS is much more than an interface bridge, in that it provides fairly complex perception and action systems to the Soar agent.

The low-level interface to ORTS is essentially a list of objects in the world, each of which has a set of attributes (like x,y position, health, owner, and size and shape information), and each of which can take any of several low-level commands (such as "move", "shoot", or "build"). The goal of our system is to build on this interface to give the Soar agent a similar interface to a human playing the game. Some of this entails adding fairly complex action controllers in the middleware, to replace what a GUI would provide a human. A person playing a commercial RTS such as StarCraft can command a unit to move to a location and build a building in one click, and can send large groups to attack the enemy semi-autonomously. Low-level behaviors like pathfinding and attack micromanagement should not be things the Soar agent is concerned with- the Soar agent should make decisions on the level of when and where to build and when and where to attack, not how to avoid an obstacle in a units path.

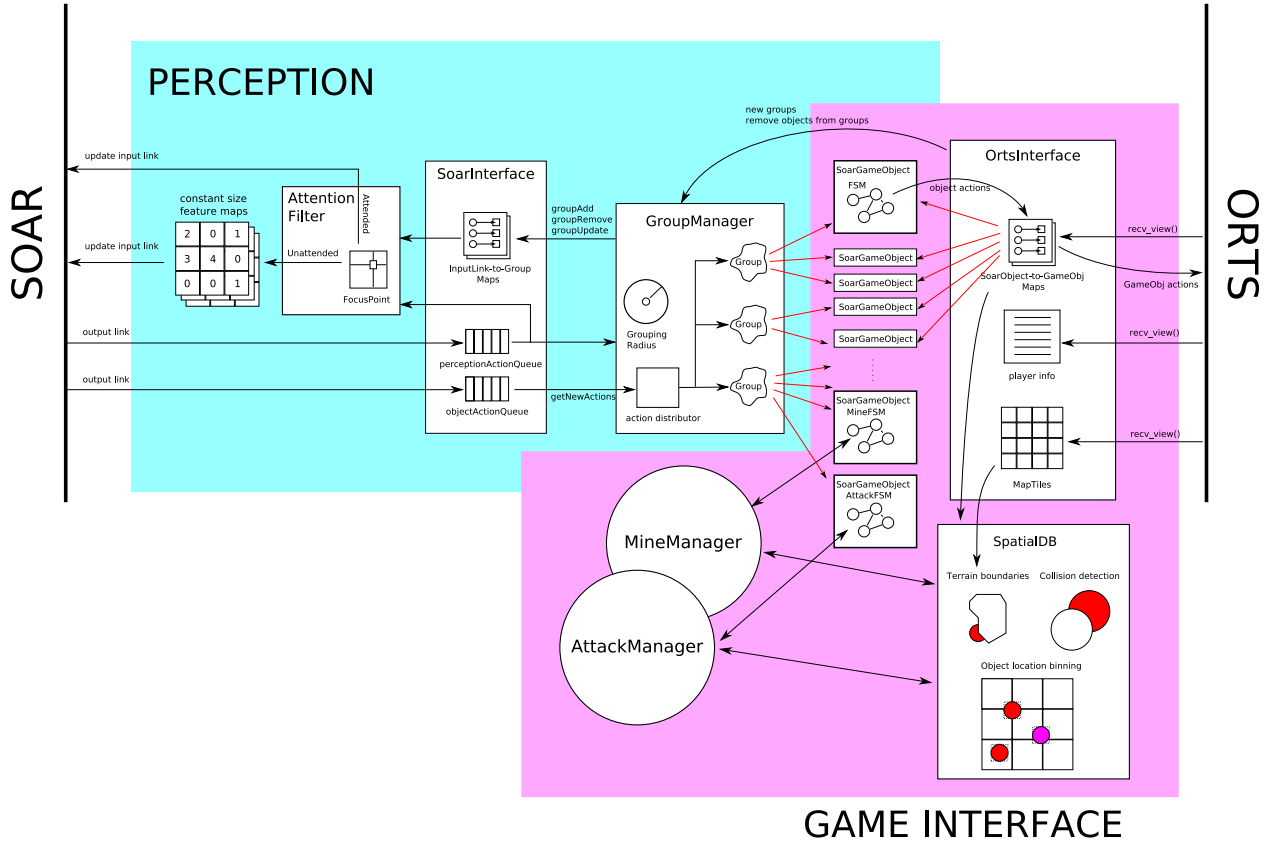


Figure 1: Overview of the modules in Sorts

Similarly, the agent's perceptions need not be on a very low level. Humans are rarely concerned with the exact x,y positions of the hundreds of units in their control, and neither should a Soar agent playing the same game. The SORTS perception system strives to give the Soar agent the same kind of information a human would receive from the GUI.

As can be seen in figure 1, SORTS is composed of several modules. The most important of these are the SoarInterface and the OrtsInterface. The SoarInterface module handles all communication between Soar and the middleware. It does this by keeping mappings between structures on the Soar input-link and the middleware-internal representation of the game state, and updating the input-link whenever the internal representation is changed. SoarInterface also keeps track of commands put on the Soar output-link and buffers them into internal action queues.

On the other hand, the OrtsInterface is the bridge between the middleware and the ORTS server. It is responsible for updating the middleware-internal game state based on the messages it receives from the ORTS server. Unlike SoarInterface, OrtsInterface does not handle translating internal action representations into actions recognizable by ORTS. That is handled directly by the behavior FSMs (see section 3).

## 2 The SORTS Perception System

During an elaborate RTS game, the amount of information available to the player is massive. Even with a "fog of war" limiting what areas of the map the player has access to, a player can see hundreds (or even thousands) of units, each with its own position, motion, and other attributes. While in general, giving a player more information is a benefit, having no filter on the raw game state would be a hindrance to the player, not only computationally (transferring too many bytes from the middleware to Soar), but also cognitively. Unfiltered information makes abstraction difficult, and makes distinctions involving locality less

apparent (for example, seeing that an enemy is within ones base).

A human RTS player has many levels of filtering available, both in the game’s GUI interface and the specialized processing of the human visual system. For example, to stage an attack, a player can restrict the GUI’s view window to his own base, and see that a group of marines is available. Perhaps one marine needs to be sent as a scout, in that case, the human player can choose to see the marine group as a set of individuals, and choose the best individual to send out. Once the individual has reached its target, the player can again observe (and select) the remaining marines as a group, and send them to attack with one mouse click. This is in contrast to the more complicated unfiltered case, where the player would see the entire map at once, search for marines, reason about whether each individual is close to the base, and perform a separate action to send each to attack.

The SORTS system has several mechanisms to enable filtering and abstraction, through the grouping of objects and focusing of visual attention. Most of the perceptions from the game world come to the agent via structures describing of groups of objects on the map. Additionally, the agent is given attentional control mechanisms to search through and select which groups in the game should be perceived. There is also a small amount of data that is presented to the agent outside of the grouping / attention system.

In addition to the descriptions here, tables 1 and 2 describe the exact input-link structure, and table 3 shows the legal attention-control commands.

## 2.1 Object Grouping

The ability of humans to see sets of similar objects as wholes has been well-studied by psychologists. This phenomenon is called Gestalt grouping ([4], [5]). The principles of Gestalt grouping specify that if objects are close together, have common features such as shape, color, and motion, they can be seen as a group. There is some top-down control of this, the observer can choose to see the individuals or the group. We use this concept to give our system the ability to perceive groups of units that are relevant to playing the game. Specifically, we group units by type, owner, and proximity. By default, groups are formed based on all three- groups are formed of units of the same type and owner which are close together. The agent can change the meaning of ”close together” by issuing a grouping-radius command to the middleware. Groups are formed by the rule that if a unit of a given class is within the grouping radius of another of the same class, they are in the same group. Note that adjusting this grouping radius to 0 will result in only one unit per group.

The agent also is able to choose to group the objects by owner alone. This is accomplished by issuing an enable-owner-grouping command to the middleware. Grouping by owner allows the agent to view the game at a higher level of abstraction, to see, for example, the locations of the different players’ bases and forces.

Actions in the game, such as *attack* or *move*, are assigned to groups. In the middleware, the execution of each unit is controlled by a separate FSM, so the behavior of the group will not necessarily be uniform. Members of a group can take different paths to a target, and so may become spatially separate as the action is executed. To simplify reasoning in the agent, all groups are automatically set to be ”sticky” once an action is assigned to them. In a sticky group, no new members can join, and members are only removed if they are killed. This way, the agent can easily see the results of an action (by checking the status of the group at a later time), and does not get confused when objects executing different actions come close together- all members of a group must be executing the same command. This is analogous to assigning a group a ”hotkey” in the user interface of a commercial RTS- the player can quickly get information on a previously-selected group, regardless of what happened to the group since it was last selected. The group will remain sticky until a *free* command is issued to it, even if all units finish executing the action.

## 2.2 Visual Attention

Even with grouping of objects, the amount of information visible can still be too much. Events in a game such as this tend to occur in restricted spatial areas, and it is worthwhile to present lots of information about a small area, while presenting less information about the rest of the world.

The human visual system provides some inspiration for a way to do this. A ”zoom lens” metaphor is often used to describe human visual attention ([3], [4], [1])– some small area of the field of vision is attended to, providing detailed information, while surrounding areas present less and less information. But this area of attention can shift in a guided manner- a human can easily jump to a red object in a sea of black, for

example. Feature Integration Theory (FIT) is a common model for describing this kind of "pop-out" effect ([6], [1]). The basic concept of FIT is that objects not attended to are not present as objects in the visual system, only as features of objects (like colors or shapes). In the red object in a field of black example, the information that something red exists (and its general area) is present whenever the object is in the visual field, but no more information about that object is known until attention selects it. Attention can be directly moved to the red object with no search, but if there were many red objects, any particular one must be found by searching all the red objects, focusing attention on each individually.

These concepts are simple to map onto our system. We simply select a certain number of groups<sup>1</sup> (adjustable by issuing *num-objects* commands) that are close to a certain point on the map, and provide all information on those groups. Then, for all groups not selected, general information is presented in the form of feature maps. A map for a given feature (enemy units, for example) consists of a list of sectors and associated unit counts. There are 9 sectors, dividing the field of view into a grid. The sectors are numbered with the upper-left as sector0 and the lower right as sector8. For each group that has a given feature, the feature count for the sector the group is centered in is incremented by the number of individuals in the group. Groups that are in attention are not shown in the feature maps.

Using the feature maps, the agent can quickly shift its attentional focus. The commands *look-at-feature* and *move-to-feature*, which take parameters of a feature name and sector number, cause the focus to jump to the center of one of the groups that shows the chosen feature in the chosen sector. This scheme, in addition to supporting very fast searches, also causes the size of the input link to be constantly bounded, preventing any problems caused by huge amounts of data overwhelming the system.

Beyond grouping and the zoom-lens area of attention, it is possible to restrict the input even further. Capability similar to shrinking a GUI window size (or dealing with a computer screen that can only show so much) is also available- completely cutting out some spatial areas from being perceived. Given our constant-bounded input size, this is not needed from a standpoint of avoiding too much data, but can serve to change the kind of data available. Specifically, since the feature maps have a set number of sectors, restricting the field of vision serves to increase their resolution. A situation where this is useful is the identification of an enemy unit in an unusual place. If the entire map is always in the field of view, a small area of it being attended to at once, the data in the feature maps will be very low resolution. If the enemy is in many places, it is likely at least one enemy unit will be somewhere in each sector. If one of those sectors also contains a region the agent is trying to control, there is no way of quickly knowing if an enemy is inside or outside the region without attending to it. However, if the agent restricts its field of view to the region in question, an enemy present in the feature map must be an invader, and can quickly be dealt with. An agent that tends to restrict its field of view to the area it is dealing with can respond quickly to unexpected objects there- they will pop out of the feature maps.

The view window is always a square, and the size is changed through issuing *change-view-width* commands to the middleware. Currently, the feature maps available are for friendly units, friendly workers, enemies, moving units, and minerals.

## 2.3 General Game Information

Although most of the preceptions given to Soar represent objects in the game directly, there is some information available that is not specific to any particular object- much of it corresponds to status text a human player would read on the screen. The agent's mineral count, a counter of elapsed time (the view frame number), along with information on the counts of various units the player owns, and some information the agent needs to know about the game configuration (player IDs and map dimensions) is all included in the game-info structure on the input link.

## 3 Low Level Control

The design philosophy of SORTS requires that Soar agents only issue high level commands comparable to those that would be expected from a human player. The reason human players of commercial RTS games

---

<sup>1</sup>Selecting a set number of groups, as opposed to a small region or a number of individuals, not only makes good sense for a video game but is supported by psychological studies [7].

do not have to worry (too much) about micromanagement is because much of it is handled for them by the game engine. Likewise, in SORTS, we handle low-level micromanagement with our middleware so that Soar does not have to. This is achieved by assigning each controllable unit a behavior specified as an finite state machine. Soar only has to specify which behavior a unit should take on, and the FSM will control the unit in predictable ways. We have already programmed a basic library of behaviors, listed in figure 2.

### 3.1 Behaviors as Finite State Machines

The main functionality of a behavior resides in its **update** function. This function is called for all living, controllable units whenever the higher level function **updateGroups()** is called in the ORTS event handler (see section 4.2). This means that the **update** function for every active behavior is called once everytime SORTS receives a viewframe from the ORTS server, which is the smallest unit of discernable state change.

During the **update** call, the behavior decides which low-level action provided by the ORTS API, or lack thereof, to execute for the game object which it controls. This decision is made based on the behavior's current state and also the state of the game object, which it has direct access to, unlike the Soar agent. The behavior also decides which state to transition to for the next **update** call.

#### 3.1.1 Coordination Managers

During preparations for the first ORTS tournament, it was decided that certain behaviors that required high levels of coordination amongst each other needed to have access to more information than just game object state if they were to perform well enough to be competitive in the tournament. Therefore, two behaviors, one that controls mining and one that controls attacking, were written to have access to a bigger picture of the overall game state via communication with coordinating mechanisms called we called "managers."

The result was that the behaviors themselves were relatively simple. During every **update** call, a managed behavior just asks the manager it was assigned to what it should do, and then does it. The advantage to using such an approach is that the manager can keep track of what each behavior it controls is doing, and coordinate sophisticated strategies such as optimal assignment of miners to mineral patches, or focusing the firepower of many units on one enemy at a time.

There are currently two such coordination managers. One of them is the mine manager, which calculates mining paths between a base and surrounding resource patches in such a way as to maximize resource income and commands the appropriate worker units to harvest those resources. The other is the attack manager, which, given a group of targets and a group of attackers, controls the attackers to eliminate the target group while trying to minimize damage done to themselves.

#### 3.1.2 Default Behaviors

There are some actions that are always advantageous to do. For example, if an offensive unit passes within firing range of an enemy unit, it is almost always better to attack the enemy unit given that there is not a more important target to attack. We don't want to burden a Soar agent with having to detect every opportunity to perform such actions, so game units can be assigned default behaviors. The desired behavior in the situation described above is implemented by the **AttackNearFSM** default behavior, which fires on nearby targets whenever they are in range and the controlled unit was not explicitly ordered to fire on something else.

As opposed to assigned behaviors, default behaviors are specified when units are created rather than when they are first desired, and a single in-game unit can have several default behaviors at once. It is important that a default behavior not override the actions of an assigned behavior. For example, the default behavior for a unit to run away when being attacked should not be manifested if the unit was assigned the move behavior in the opposite direction of the retreat. However, there is no built-in mechanism to ensure this kind of suppression in SORTS yet, so it is the responsibility of the default behavior to make sure it doesn't obstruct assigned behaviors.

Behavior name	Parameters	Description
AttackFSM	id	Attack group with specified id
AttackNearFSM		Attack any enemy target within range, except when already attacking another target
BuildFSM	t, x, y	Build a building of type t at (x, y)
MineFSM		Mine for minerals
MoveFSM	x, y, speed	Object tries to pathfind and move to coordinate (x, y)
TrainFSM	t	Train a unit of type t

Figure 2: A list of implemented behaviors.

## 4 Updating the Game State

### 4.1 Raw State Information

The ORTS API provides complete game state update information in the form of 6 lists:

- **new\_tile\_indexes** New game tiles encountered via uncovering of fog of war.
- **new\_objs** New game objects, such as enemy units and buildings, or world objects such as trees.
- **changed\_objs** Game objects that have changed since the previous viewframe.
- **vanished\_objs** Game objects that disappeared from the player’s vision.
- **dead\_objs** Game objects that died.
- **new\_boundaries** New terrain boundaries that demarcate the border between different terrain types, such as ground and cliff.

SORTS currently takes into account the information in all the lists except **new\_tile\_indexes** and **new\_boundaries**. The former is ignored because SORTS currently does not have any functioning capability to reason about terrain, except for collision detection. The latter is ignored because we get this information from the game tiles directly whenever we check for boundaries. Note also that currently, SORTS treats objects vanishing (reported in **vanished\_objs**) and objects dying (reported in **dead\_objs**) as one and the same thing. This will probably change in the future as agents use more sophisticated reasoning.

All game state updates, both internal to the middleware and to the Soar input link, are performed in the ORTS event handler, which is triggered everytime the middleware receives a new viewframe from the ORTS server. An overview of the flow of control among the various functions is shown in figure 3. Detailed descriptions of the two event handlers are provided below.

### 4.2 The ORTS Event Handler

The ORTS event handler is triggered everytime the middleware receives a new viewframe from the ORTS server. This event handler and the Soar event handler are the main functions from which most other function calls are made. Pseudo-code for the handler is shown in figure 4.

Each function call is described below.

- **lockMutex()** and **unlockMutex()** This mutex locks the Soar event handler from executing until the ORTS event handler is finished. This is important since we don’t want the Soar command buffer to change while we are processing that buffer.
- **mergeChanges(oldChanges, changes)** If SORTS falls behind the ORTS server too much, changes to the game state will be accumulated but not processed in the interest of catching up. This function merges changes reported in previous viewframes with the current changes.

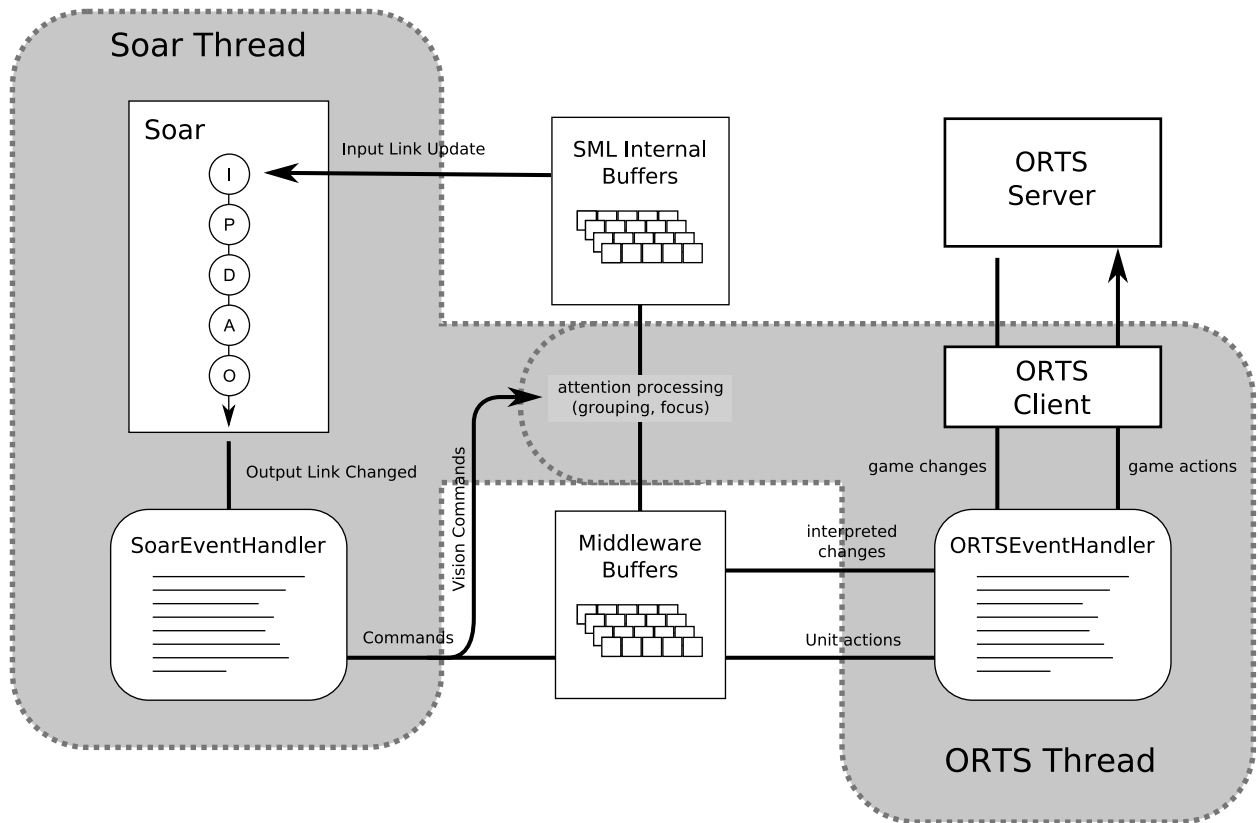


Figure 3: Overview of how control is passed through Sorts. Note that the two gray regions represent two different threads executing asynchronously.

```

ORTSEventHandler
  lockMutex()
  mergeChanges(oldChanges, changes)
  if currentViewFrame - lastActionFrame > ALLOWED_LAG then
    unlockMutex()
    return
  end
  removeDeadObjects(changes)
  assignSoarActions()
  updateSoarGameObjects(changes)
  sendActions()
  updateGroups()
  unlockMutex()
end

```

Figure 4: Pseudo code for the ORTS event handler

```

SoarEventHandler
  lockMutex()
  if Catchup = true then
    unlockMutex()
    return
  end if
  getNewSoarOutput()
  processVisionCommands()
  processGameCommands()
  unlockMutex()
end

```

Figure 5: Pseudo code for the Soar event handler

- **removeDeadObjects(changes)** This function removes all game objects that have just died or vanished in the current frame. This call needs to be made before the call to **assignSoarActions()** so that we don't try to assign any actions to objects that disappeared, which the ORTS server would not understand.
- **assignSoarActions()** Interprets commands queued from the Soar output-link as ORTS actions and distributes them to the appropriate groups. The actions are not sent to the ORTS server until **sendActions()** is called.
- **updateSoarGameObjects(changes)** This is the main function that updates the attributes of the SoarGameObject structures in the middleware. It is called after **assignSoarActions()** so that units will begin responding to their commands in the Soar cycle that immediately follows the one in which they were assigned.
- **sendActions()** This is a call to the ORTS API that sends all queued actions to the server.
- **updateGroups()** Runs the grouping algorithm over SoarGameObjects that have changed or appeared, and prunes those groups whose members have died. These changes are directly reflected in the Soar input-link, but are buffered in the SML interface until the next Soar decision cycle.

### 4.3 The Soar Event Handler

The Soar event handler is triggered at the end of each Soar decision cycle. The main responsibility of this function is to take commands off the Soar output-link and buffer them into action queues in the middleware.

Figure 5 shows the pseudo-code for the function. The following is a list of descriptions for each call made in the function.

- **lockMutex()** and **unlockMutex()** These calls lock on the same mutex as the ORTS event handler, thereby making the execution of these two functions mutually exclusive.
- **getNewSoarOutput()** Takes all commands on the Soar output-link and queues them into lists in the middleware. This includes both commands issued to in-game units as well as commands that adjust middleware parameters such as grouping radius and center of visual attention. However, none of the commands are actually processed by this function.
- **processVisionCommands()** Processes Soar commands related to the perceptual system. These commands include changing the grouping radius, looking at a specific coordinate, looking at a feature in the feature map, and changing the maximum number of objects allowed on the input-link at any time (for a complete list, see section 5.2. These commands are processed here rather than with the in-game commands because Soar decision cycles can occur much faster than ORTS events.



- `processGameCommands()` Handles all commands that affect the way the middleware executes low-level actions, such as in the FSMs, but do not translate directly into ORTS game object actions. Also handles miscellaneous queries that Soar may make to the middleware for information not normally provided to it. Currently, there are only three such commands: finding a location for a building, and setting and clearing the mineral buffer (see section 5.2).

## 5 Soar IO Description

### 5.1 The SORTS input-link

There are five top-level attributes on the SORTS input link, "groups", "game-info", "feature-maps", "vision-info", and "query-results". The groups, feature-maps, and vision-info structures are all part of the main visual system (see section 2), while game-info contains higher-level information about the game world, and query-results is used to communicate the results of specialized queries from Soar to the middleware.

The exact data structures of the input-link are located in tables 1 and 2.

### 5.2 The SORTS output-link

The output-link allows the Soar agent to act in the game world by issuing commands to groups of friendly units, communicate directly to the middleware for actions such as querying, and adjust the parameters of the visual subsystem in the middleware. All structures on the output link are similar- the Soar agent creates an `output-link.command` object with a name (**output-link.command.name**) specified. Depending on the command name, different parameters are needed. Parameters are attached to the command structure, as in **output-link.command.param0**. All legal commands and their necessary parameters are detailed in table 3.

## References

- [1] John R. Anderson. *Cognitive Psychology and its Implications*. Worth Publishers, New York, 2004.
- [2] Michael Buro. Orts: A hack-free rts game environment. In *Proceedings of the International Computers and Games Conference 2002*, 2002.
- [3] St James JD. Eriksen CW. Visual attention within and around the field of focal attention: a zoom lens model. *Percept Psychophys.*, 40:225–40, Oct. 1986.
- [4] R. Hill. Modeling perceptual attention in virtual humans. In *Proc. 8 th Conf. Computer Generated Forces and Behavioral Representation*, 1999.
- [5] Michael Kubovy, Alex O. Holcombe, and Johan Wagemans. On the lawfulness of grouping by proximity,. *Cognitive Psychology*, 35:71–98, Feb 1998.
- [6] Philip T. Quinlan. Visual feature integration theory : Past, present, and future. *Psychological Bulletin*, 129:643–673, Sept. 2003.
- [7] Brian J. Scholl. Objects and attention: the state of the art. *Cognition*, 80:1–46, Jun 2001.

Attributes of io.input-link	
attribute	description
<b>vision-info structure:</b>	
vision-info	Contains information on the current state of the vision system.
vision-info.center-x vision-info.center-y	The coordinate of the center of the region in view.
vision-info.focus-x vision-info.focus-y	The coordinate of the center of focus (spotlight of attention).
vision-info.num-objects-visible	The maximum number of objects (groups) present on the input-link. All other objects within the view window are present in feature maps.
vision-info.grouping-radius	All objects of the same type (except as below) and owner within this distance of each other are in the same group (set to 0 for individuals).
vision-info.owner-grouping	Ignore type when grouping, only group by owner (1 if enabled, 0 if disabled).
<b>groups structure:</b>	
groups	The set of groups being attended to.
groups.group	Multi-valued, one instance for each group. Detailed in table 2.
<b>feature-maps structure:</b>	
feature-maps	Contains all feature maps- low-resolution information about certain features of unattended (but visible) objects.
feature-maps.friendly	Friendly feature map- each friendly unit (not group) results in one instance of this feature, marked in the sector of the group's center of gravity.
feature-maps.friendly-workers	A subset of the friendly feature map, showing only workers.
feature-maps.enemy	Similar to the above, but for enemy units.
feature-maps.minerals	Similar to the above, but for minerals.
feature-maps.moving-units	Moving units (if grouping is used, one moving object causes the whole group to be seen as moving).
feature-maps.(any).sector0 .. feature-maps.(any).sector8	Feature counts for each of the nine sectors. 0 is the upper left, 8 is the lower right.
<b>game-info structure:</b>	
game-info	General information about the game (non-visual information).
game-info.num-players	The number of players.
game-info.player-id	The ID number of the Soar player.
game-info.map-xdim game-info.map-ydim	The dimensions of the game map.
game-info.view-frame	The last view frame handled by the middleware (the number of game cycles executed).
game-info.my-minerals	Minerals available to the player.
game-info.worker-count game-info.marine-count game-info.tank-count	The number of units of each type owned by the player.
<b>query-results structure:</b>	
query-results	The result of the last query to the middleware.
query-results.query-name	The name of the last query executed.
query-results.param0 query-resutls.param1	Query return values- the meaning of these is dependant on the query-name.

Table 1: The SORTS input-link

Attributes of io.input-link.groups.group objects		
attribute	which groups	description
num-members	all	The number of individuals comprising the group.
type	all	The type of the group (ex: worker, mineral).
x-pos y-pos	all	The x,y location of the center of gravity of the group.
x-min x-max y-min y-max	all	The bounding box of the group.
health	all	The sum of the health of all units in the group.
taking-damage	all	The number of members of the group currently taking damage (under attack).
shooting	all	The number of members of the group currently attacking an enemy.
speed	all	The average speed of the group.
heading	all	The average heading of the group.
dist-to-focus	all	The distance from the center of gravity of the group to the attentional focus point.
dist-to-query	all	The distance from the center of gravity of the group to the last query location.
owner	all	The player number of the group's owner.
enemy	all	1 if the group belongs to an enemy player, 0 otherwise.
sticky	friendly	1 if the group is sticky- sticky groups remain together even if they are no longer spatially close.
command	friendly	The last command issued to the group ("none" if no command has been issued).
command-running	friendly	The number of members of the group currently executing a command.
command-success	friendly	The number of members of the group that successfully completed the last command.
command-failure	friendly	The number of members of the group that unsuccessfully completed the last command.
minerals	friendly workers	The total number of minerals possessed by the workers in the group.
active-mining	friendly workers	The number of workers that are actively mining.

Table 2: Detail of group objects.

Soar output-link commands		
command name	parameters	description
<b>Vision commands</b>		
grouping-radius	<i>value</i>	Change the grouping radius of the vision system.
enable-owner-grouping	(none)	Enable grouping-by-owner.
disable-owner-grouping	(none)	Disable grouping-by-owner.
change-view-width	<i>value</i>	Change the width of the agent's view to be <i>value</i> .
look-at-location	<i>x, y</i>	Move the focus of attention to the given coordinate (which must be in the current view window).
move-to-location	<i>x, y</i>	Shift the view window to be centered at the given coordinate, and make that the focus of attention.
look-at-feature	<i>feature, sector</i>	Move the focus of attention to a given feature in a given sector. The legal features are the names of the feature-map objects on the input link.
move-to-feature	<i>feature, sector</i>	As above, but re-center the view window to the new location, also.
num-objects	<i>value</i>	Change the maximum number of objects (groups) present on the input-link at once.
<b>General middleware commands</b>		
locate-building	<i>building, x, y, distance</i>	This command requests that the middleware attempt to find a location for a building of type <i>building</i> approximately <i>distance</i> units from the coordinate given. The resulting coordinate is returned through the query-result structure on the input-link.
increase-mineral-buffer	<i>value</i>	Increase the mineral buffer by <i>value</i> minerals.
clear-mineral-buffer	(none)	Set the mineral buffer to 0.
<b>Group commands</b>		
move	<i>group0, param0, param1</i>	Move the group with ID <i>group0</i> to the x,y location where <i>param0</i> is x and <i>param1</i> is y, using the default precision of 10 (allow units to complete successfully if they are within 10 of the target).
move	<i>group0, param0, param1, param2</i>	Move group <i>group0</i> to the x,y location where <i>param0</i> is x and <i>param1</i> is y, using a precision of <i>param2</i> .
build	<i>group0, param0..param3</i>	Use group <i>group0</i> to build a building of type <i>param0</i> <sup>a</sup> at the x,y location given by <i>param1, param2</i> . If <i>param3</i> is 1, the mineral buffer is used, otherwise it is not.
train	<i>group0, param0..param2</i>	Use group <i>group0</i> (a building) to train units of type <i>param1</i> <sup>b</sup> . Train up to <i>param2</i> units, and use the mineral buffer if <i>param3</i> is 1.
attack	<i>group0, group1</i>	Use the friendly group with ID <i>group0</i> to attack <i>group1</i> .
mine	<i>group0</i>	Assign <i>group0</i> to mine minerals. The control center and mineral patch to use are automatically determined in the middleware.
stick	<i>group0</i>	Set <i>group0</i> 's to be sticky- ensure it stays as one group, even if the members move apart. Assigning any of the above actions to a group does this by default.
free	<i>group0</i>	Clear <i>group0</i> 's sticky status- allow the members to be split up and join other groups.
join	<i>group0, group1</i>	Force <i>group0</i> 's members to join <i>group1</i> . If <i>group1</i> is not sticky, this may be automatically undone the next cycle.

<sup>a</sup>Legal building types: 0=controlCenter, 1=barracks, 2=factory

<sup>b</sup>Legal unit types: 0=worker, 1=marine, 2=tank

Table 3: SORTS output-link commands- see section 5.2 for details on common command structure.